

Papers mentioned in week 5

Here's a list of papers mentioned in week 5 slides.

Video 1:

- A. Bhargava, A. Celikyilmaz, D. Hakkani-Tur, and R. Sarikaya. **EASY CONTEXTUAL INTENT PREDICTION AND SLOT DETECTION** (2013). <http://www.cs.toronto.edu/~aditya/publications/contextual.pdf>
- K. Scheffler and S. Young. **Simulation of human-machine dialogues** (1999). <http://mi.eng.cam.ac.uk/~sjy/papers/scyo99.ps.gz>

Video 2:

- Wenpeng Yin, Katharina Kann, Mo Yu, Hinrich Schütze. **Comparative Study of CNN and RNN for Natural Language Processing** (2017). <https://arxiv.org/pdf/1702.01923.pdf>
- Yann N. Dauphin, Angela Fan, Michael Auli, David Grangier. **Language Modeling with Gated Convolutional Networks** (2017). <https://arxiv.org/pdf/1612.08083.pdf>
- Jonas Gehring, Michael Auli, David Grangier, Denis Yarats, Yann N. Dauphin. **Convolutional Sequence to Sequence Learning** (2017). <https://arxiv.org/pdf/1705.03122.pdf>
- Bing Liu, Ian Lane. **Attention-Based Recurrent Neural Network Models for Joint Intent Detection and Slot Filling** (2016). https://www.isca-speech.org/archive/Interspeech_2016/pdfs/1352.PDF
- Gokhan Tur, Dilek Hakkani-Tur, Larry Heck. **WHAT IS LEFT TO BE UNDERSTOOD IN ATIS?** (2010). <https://www.microsoft.com/en-us/research/wp-content/uploads/2010/12/SLT10.pdf>

Video 3:

- Yun-Nung Chen, Dilek Hakkani-Tur, Gokhan Tur, Jianfeng Gao, and Li Deng. **End-to-End Memory Networks with Knowledge Carryover for Multi-Turn Spoken Language Understanding** (2016). https://www.microsoft.com/en-us/research/wp-content/uploads/2016/06/IS16_ContextualSLU.pdf

Video 4:

- Jason P.C. Chiu, Eric Nichols. **Named Entity Recognition with Bidirectional LSTM-CNNs** (2016). <https://arxiv.org/pdf/1511.08308v4.pdf>

Video 5:

- Xijun Li, Yun-Nung Chen, Lihong Li, Jianfeng Gao, Asli Celikyilmaz. **Investigation of Language Understanding Impact for Reinforcement Learning Based Dialogue Systems** (2017). <https://arxiv.org/pdf/1703.07055.pdf>
- Matthew Henderson, Blaise Thomson and Jason Williams. **Dialog State Tracking Challenge 2 & 3** (2013). <http://camdial.org/~mh521/dstc/downloads/handbook.pdf>
- Nikola Mrksic, Diarmuid O Seaghdha, Tsung-Hsien Wen, Blaise Thomson, Steve Young. **Neural Belief Tracker: Data-Driven Dialogue State Tracking** (2017). <https://arxiv.org/pdf/1606.03777.pdf>
- Layla El Asri, et al. **FRAMES: A CORPUS FOR ADDING MEMORY TO GOAL-ORIENTED DIALOGUE SYSTEMS** (2017). <https://arxiv.org/pdf/1704.00057.pdf>

Video 6:

- Xuesong Yang, Yun-Nung Chen, Dilek Hakkani-Tur, Paul Crook, Xijun Li, Jianfeng Gao, Li Deng. **END-TO-END JOINT LEARNING OF NATURAL LANGUAGE UNDERSTANDING AND DIALOGUE MANAGER** (2017). <https://arxiv.org/pdf/1612.00913.pdf>

Keep up-to-date with NLP research

Congratulations for finishing with the lectures in our course!

NLP is a huge and rapidly emerging area. So to have an up-to-date understanding of its advances one should always keep track of what is going. In these reading material we provide some links for you that give a nice **overview of NLP trends as for the end of 2017**.

First, it is always a good idea to check out highlights from main conferences. There are nicely summarized trends of ACL-2017: [part 1](#), [part 2](#). Also, some highlights from EMNLP-2017 are available [here](#). Second, it would be a good idea to monitor some blogs, e.g. Sebastian Ruder has nice posts about [DL in NLP](#), [optimization trends](#), [word embeddings](#), and many others.

One of still active topics is Thought Vectors and how one can interpret directions in the hidden space. E.g. you might be interested to check out [this post](#). However, it's getting more clear that compressing all the input into one vector is often not enough and one might make nice things with [attention and linguistic information](#). Some more tips about attention [here](#).

Finally, [this is](#) another nice overview of 2017 trends in NLP research - advances in unsupervised machine translation seem especially exciting!

Not surprisingly, you will notice that each new year introduces new SOTA models and NLP techniques. Just to mention a few:

- Transformers are everywhere, [this blogpost](#) might be a good place to start
- New representation learning techniques (BERT, ELMO, etc).: [blogpost](#)
- Big pre-trained LM models (Transformer-XL, GPT-2, etc): [blogpost](#)

To conclude, we would like to say **thank you** for taking our course and wish best of luck in your future NLP projects!